Size Balanced Tree

From PEGWiki

A **size balanced tree** (**SBT**) is a self-balancing binary search tree (BBST) first published by Chinese student Qifeng Chen in 2007. The tree is rebalanced by examining the sizes of each node's subtrees. Its abbreviation resulted in many nicknames given by Chinese informatics competitors, including "sha bi" tree (Chinese: 傻屄树; pinyin: *shǎ bī shù*; literally meaning "dumb cunt tree") and "super BT", which is a homophone to the Chinese term for snot (Chinese: 鼻涕; pinyin: *bíti*) suggesting that it is messy to implement. Contrary to what its nicknames suggest, this data structure can be very useful, and is also known to be easy to implement. Since the only extra piece of information that needs to be stored is sizes of the nodes (instead of other "useless" fields such as randomized weights in treaps or colours in red–black tress), this makes it very convenient to implement the *select-by-rank* and *get-rank* operations (easily transforming it into an order statistic tree). It supports standard binary search tree operations such as insertion, deletion, and searching in O(log *n*) time. According to Chen's paper, "it works much faster than many other famous BSTs due to the tendency of a perfect BST in practice."[1]

Contents

Properties

The size balanced tree examines each node's size (i.e. the number of nodes below it in the subtree) to determine when rotations should be performed. Each node T in the tree satisfies the following two properties:

1. $size(T.left) \geq size(T.right.left), size(T.right.right)$
2. $size(T.right) \geq size(T.left.left), size(T.left.right)$

In other words, each child node of T is not smaller in size than the child nodes of its sibling. Clearly, we should consider the sizes of nonexistent children and siblings to be 0.

Consider the following example where T is the node in question, L, R are its child nodes, and A, B, C, D are subtrees which also satisfy the above SBT properties on their own.

```
      T
     / \
    /   \
   L     R
  / \   / \
 A   B C   D
```

Then, the node T must satisfy:

- $size(L) \geq size(C), size(D)$
- $size(R) \geq size(A), size(B)$

Rotations

The rotations of SBTs are analogous to those in other self-balancing binary search trees.

```
 -------------       Right Rotation      -----------
|    Q    |        --------------->      |    P    |
|   / \   |                              |   / \   |
 -- P   C |                              |  A   Q --
   / \       <---    Left Rotation    --->   / \
  A   B      <--------------                B   C
```

Left Rotation

```
def left-rotate(t):
    k ← t.right
    t.right ← k.left
    k.left ← t
    k.size ← t.size
    t.size ← t.left.size + t.right.size + 1
    t ← k
```

Right Rotation

```
def right-rotate(t):
    k ← t.left
    t.left ← k.right
    k.right ← t
    k.size ← t.size
    t.size ← t.left.size + t.right.size + 1
    t ← k
```
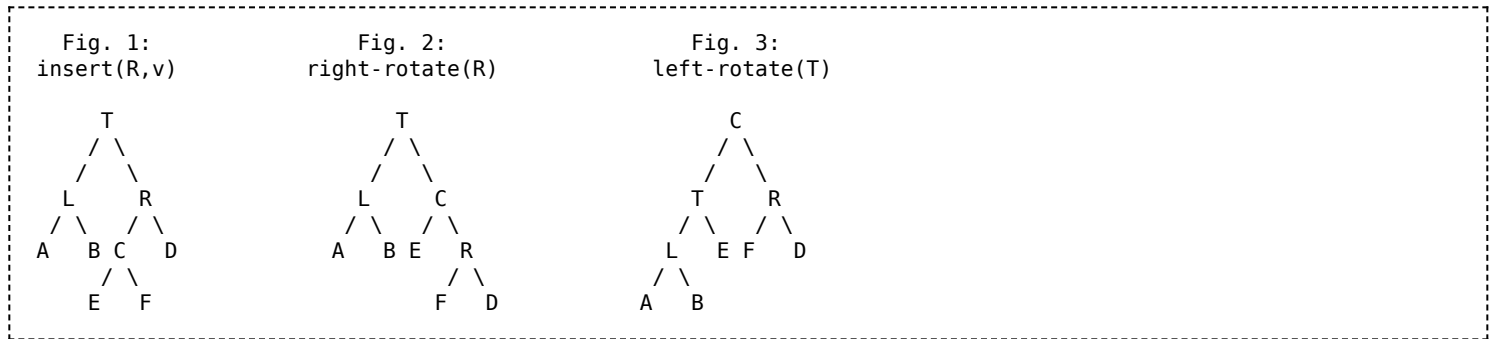
Maintenance

After insertions and deletions, the new sizes of subtrees may violate the two properties above. Thus, we define a procedure maintain(T) to rebalance the SBT rooted at the node T. This should be called with the precondition that T's children are already SBTs themselves. Since property 1 and 2 are symmetrical, we will only discuss property 1.

There are 4 cases to consider when rebalancing.

- **Case 1**: $size(T.left) < size(T.right.left)$

  Perhaps after inserting a value to $T.right$, the scenario below (figure 1) may occur, leading to $size(L) < size(C)$.
  To fix this, we first perform a `right-rotate` on $T.right$ (figure 2) and then a `left-rotate` on T (figure 3).

```
  Fig. 1:              Fig. 2:               Fig. 3:
 insert(R,v)         right-rotate(R)        left-rotate(T)

    T                    T                     C
   / \                  / \                   / \
  /   \                /   \                 /   \
 L     R              L     C               T     R
/ \   / \            / \   / \             / \   / \
A  B C   D           A  B E   R           L   E F   D
     / \                    / \          / \
    E   F                  F   D         A   B
```

  After these operations, the properties of the entire tree in figure 3 becomes unpredictable. Luckily, the subtrees A, B, D, E, F, L are still SBTs. Thus, we can recursively call `maintain` on subtrees R and T to take care of them.
  Now that all of the subtrees are SBTs, we still have to make sure that the root node C satisfies the SBT properties. So, we call `maintain` one last time on root node C.

- **Case 2**: $size(T.left) < size(T.right.right)$

  Perhaps after inserting a value to $T.right$, the scenario below (figure 4) may occur, leading to $size(L) < size(D)$. This is similar to case 1, except that instead of going below C, E and F instead goes below D. We can omit them from the diagram.
  Fixing this, we will perform a `left-rotate` on the root node T, obtaining the structure in figure 5.

```
  Fig. 4:              Fig. 5:
 insert(R,v)         left-rotate(T)
    T                     R
   / \                   / \
  /   \                 /   \
 L     R               T     D
/ \   / \             / \
A  B C   D           L   C
                    / \
                   A   B
```

  After this, the tree rooted at R is still not yet a SBT because $size(C) < size(A)$ or $size(C) < size(B)$ may be true. So, we continue to call `maintain` on T.
  Now that we have satisfied the precondition of making R's subtrees SBTs, we may call `maintain` on R itself.

- **Case 3**: $size(T.right) < size(T.left.right)$

  Symmetrical to case 1.

- **Case 4**: $size(T.right) < size(T.left.left)$

  Symmetrical to case 2.

With this casework being taken care of, it becomes straightforward to actually implement `maintain`.

```
def maintain(t):

    if t.left.size < t.right.left.size:          //case 1
        right-rotate(t.right)
        left-rotate(t)
        maintain(t.left)
        maintain(t.right)
        maintain(t)

    else if t.left.size < t.right.right.size:    //case 2
        left-rotate(t)
        maintain(t.left)
        maintain(t)

    else if t.right.size < t.left.right.size:    //case 1'
        left-rotate(t.left)
        right-rotate(t)
        maintain(t.left)
        maintain(t.right)
        maintain(t)

    else if t.right.size < t.left.left.size:     //case 2'
        right-rotate(t)
        maintain(t.right)
        maintain(t)
```

This pseudocode is slightly slow and redundant. Since we know that the two SBT properties will *usually* be satisfied, the following is an optimization. Simply add an extra boolean flag to the `maintain` parameters, indicating whether cases 1/2 or their symmetrical cases are being examined. If the flag is TRUE, then we examine cases 1 and 2, otherwise we examine cases 3 and 4. Doing so will eliminate many unnecessary comparisons.

```
def maintain(t, flag):

    if flag:
        if t.left.size < t.right.left.size:          //case 1
            right-rotate(t.right)
            left-rotate(t)
        else if t.left.size < t.right.right.size:    //case 2
            left-rotate(t)
        else:
            done
    else:
        if t.right.size < t.left.right.size:         //case 1'
            left-rotate(t.left)
            right-rotate(t)
        else if t.right.size < t.left.left.size:     //case 2'
            right-rotate(t)
        else:
            done

    maintain(t.left, FALSE)    //maintain the left subtree
    maintain(t.right, TRUE)    //maintain the right subtree
    maintain(t, TRUE)          //maintain the whole tree
    maintain(t, FALSE)         //maintain the whole tree
```

The proof for why `maintain(t.left, TRUE)` and `maintain(t.right, FALSE)` are unnecessary can be found in section 6 of Chen's paper. Furthermore, the running time of `maintain` is O(1) amortized (which means that you do not have to worry about it not terminating).

Fundamental Operations

Searching

Searching in SBTs is exactly the same as searching in other binary search trees. The following iterative implementation will return a pointer to the node in the SBT rooted at t which has key k.

```
def search(t, k):
    x ← t
    while x is not NIL:
        if k < x.key then x ← x.left
        else if x.key < k then x ← x.right
        else return x
    return NIL    //key not found!
```

Get Max/Min

The size of the SBT is already stored. These operations can thus be handled trivially by the `select` operation implemented in the section below.

Iteration

Iterating a SBT is exactly the same as iterating a normal binary search tree (by repeatedly finding nodes' predecessors/successors).

Insertion

Inserting into a SBT is very simple. The only difference from normal binary search trees is that it has an extra call to `maintain` at the end. The following recursive version will insert the node x into the SBT rooted at t.

```
def insert(t, x):
    if x is NIL:
        t ← x
    else
        t.size ← t.size + 1
        if x.key < t.key:
            insert(t.left, x)
        else
            insert(t.right, x)
    maintain(t, x.key ≥ t.key)
```

Deletion

Deletion is exactly the same as in normal binary search trees. It is not even necessary to call `maintain` afterwards. The proof for this is as follows: A SBT will have all of its properties before deletion. Even though we cannot guarantee that the SBT will retain its balanced properties after the insertion, we know for sure that its height (and thus, its running time) will not increase. Given this, it is clear that calling `maintain` after deleting is extraneous.

Order Statistics

Since SBTs already conveniently store the $size$ field to maintain balance, nothing else is needed to transform it into a fully-fledged order statistics tree.

Select

The following function returns a pointer to the ith smallest element in the SBT rooted at t, where i is zero-indexed. To make this one-indexed, simply change "`r ← t.left.size`" to "`r ← t.left.size + 1`" and "`i - (r + 1)`" to "`i - r`".

```
def select(t, i):
    r ← t.left.size
    if i = r:
        return t
    else if i < r:
        return select(t.left, i)
    else
        return select(t.right, i - (r + 1))
```

Rank

Determining the rank of an element in a SBT is exactly the same as doing so for a regular binary search tree.

References

1. Chen, Qifeng. "Size Balanced Tree" (http://www.scribd.com/doc/3072015/), Guandong, China, 29 December 2006.

Retrieved from "http://wcipeg.com/wiki/index.php?title=Size_Balanced_Tree&oldid=1819"