

Part 3. The neural network paradigm

We're seeing a huge push from industries to place neural networks on a pedestal. Deep-learning research has become a corporate status symbol, with the theory behind it obfuscated by smoke and mirrors. Massive amounts of money have been thrown at marketing this technology by companies including NVIDIA, Facebook, Amazon, Microsoft, and, let's not forget, Google. Regardless, deep learning works exceptionally well for solving some problems, and using Tensor-Flow is how we'll implement it.

The chapters in this part of the book introduce neural networks from the basics and apply these architectures to real-world practical applications. In order, the chapters are about autoencoders, reinforcement learning, convolutional neural networks, recurrent neural networks, sequence-to-sequence models, and ranking. Full speed ahead!

[Recommended](#) / [Playlists](#) / [History](#) / [Topics](#) / [Tutorials](#) / [Settings](#) / [Get the App](#) / [Sign Out](#)